

Predicting Sale Prices Using Advanced Regression and Feature Engineering

In this notebook I have explored the House Prices dataset from Kaggle using several regression algorithms. We'll go through the complete machine learning workflow step by step:

- 📊 Data exploration & cleaning
- 🌐 Feature engineering
- 🧠 Model building & tuning
- 🎯 Model blending (XGBoost, LightGBM, Ridge, Lasso)
- 📈 Generating final predictions

The end goal is to achieve a competitive relative error in prediction of the Saleprice of Houses with maximum accuracy.

We begin by importing essential libraries for data handling, visualization, and modeling. This includes Pandas, NumPy, Seaborn, Matplotlib, and multiple machine learning libraries like Scikit-learn, XGBoost, LightGBM, and CatBoost.

```
In [36]: # Ignore warnings
import warnings
warnings.filterwarnings('ignore')

# Data handling
import numpy as np
import pandas as pd

# Visualization
import matplotlib.pyplot as plt
import seaborn as sns

# Modeling
from sklearn.model_selection import train_test_split, cross_val_score, KFold
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.ensemble import GradientBoostingRegressor, RandomForestRegressor
import xgboost as xgb
import lightgbm as lgb
import catboost as cb

# Display settings
pd.set_option('display.float_format', lambda x: '%.3f' % x)
```

```
# Load data
train = pd.read_csv('/kaggle/input/house-prices-advanced-regression-techni
test = pd.read_csv('/kaggle/input/house-prices-advanced-regression-techni

print("Train shape:", train.shape)
print("Test shape:", test.shape)
train.head()
```

Train shape: (1460, 81)

Test shape: (1459, 80)

Out[36]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	Lan
0	1	60	RL	65.000	8450	Pave	NaN	Reg	
1	2	20	RL	80.000	9600	Pave	NaN	Reg	
2	3	60	RL	68.000	11250	Pave	NaN	IR1	
3	4	70	RL	60.000	9550	Pave	NaN	IR1	
4	5	60	RL	84.000	14260	Pave	NaN	IR1	

5 rows × 81 columns

Key **observations** after loading the datasets:

- Train set has 1460 rows and 81 columns
- Test set has 1459 rows and 80 columns
- SalePrice is our target variable, available only in the training set.

Initial Data Exploration

We examine the dataset to understand its structure and identify missing values.

- Many categorical columns like `PoolQC`, `Fence`, `Alley`, `MiscFeature` have large sections of missing data.
- Numeric features like `LotFrontage`, `GarageYrBlt`, and `MasVnrArea` have moderate missing values.

In [37]:

```
# Overview of training data
print("---- TRAIN DATA INFO ----")
train.info()

print("\n---- TEST DATA INFO ----")
test.info()
```

----- TRAIN DATA INFO -----

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 81 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Id               1460 non-null    int64  
 1   MSSubClass        1460 non-null    int64  
 2   MSZoning          1460 non-null    object  
 3   LotFrontage       1201 non-null    float64 
 4   LotArea           1460 non-null    int64  
 5   Street            1460 non-null    object  
 6   Alley              91 non-null     object  
 7   LotShape           1460 non-null    object  
 8   LandContour        1460 non-null    object  
 9   Utilities          1460 non-null    object  
 10  LotConfig          1460 non-null    object  
 11  LandSlope          1460 non-null    object  
 12  Neighborhood       1460 non-null    object  
 13  Condition1         1460 non-null    object  
 14  Condition2         1460 non-null    object  
 15  BldgType           1460 non-null    object  
 16  HouseStyle          1460 non-null    object  
 17  OverallQual        1460 non-null    int64  
 18  OverallCond         1460 non-null    int64  
 19  YearBuilt           1460 non-null    int64  
 20  YearRemodAdd        1460 non-null    int64  
 21  RoofStyle           1460 non-null    object  
 22  RoofMatl            1460 non-null    object  
 23  Exterior1st          1460 non-null    object  
 24  Exterior2nd          1460 non-null    object  
 25  MasVnrType          588 non-null     object  
 26  MasVnrArea          1452 non-null    float64 
 27  ExterQual            1460 non-null    object  
 28  ExterCond            1460 non-null    object  
 29  Foundation           1460 non-null    object  
 30  BsmtQual            1423 non-null    object  
 31  BsmtCond            1423 non-null    object  
 32  BsmtExposure         1422 non-null    object  
 33  BsmtFinType1         1423 non-null    object  
 34  BsmtFinSF1           1460 non-null    int64  
 35  BsmtFinType2         1422 non-null    object  
 36  BsmtFinSF2           1460 non-null    int64  
 37  BsmtUnfSF            1460 non-null    int64  
 38  TotalBsmtSF          1460 non-null    int64  
 39  Heating              1460 non-null    object  
 40  HeatingQC             1460 non-null    object  
 41  CentralAir            1460 non-null    object  
 42  Electrical            1459 non-null    object  
 43  1stFlrSF              1460 non-null    int64  
 44  2ndFlrSF              1460 non-null    int64  
 45  LowQualFinSF          1460 non-null    int64  
 46  GrLivArea             1460 non-null    int64  
 47  BsmtFullBath          1460 non-null    int64  
 48  BsmtHalfBath          1460 non-null    int64  
 49  FullBath              1460 non-null    int64  
 50  HalfBath              1460 non-null    int64  
 51  BedroomAbvGr          1460 non-null    int64  
 52  KitchenAbvGr          1460 non-null    int64  
 53  KitchenQual            1460 non-null    object
```

```

54 TotRmsAbvGrd    1460 non-null    int64
55 Functional     1460 non-null    object
56 Fireplaces      1460 non-null    int64
57 FireplaceQu    770 non-null     object
58 GarageType      1379 non-null    object
59 GarageYrBlt    1379 non-null    float64
60 GarageFinish    1379 non-null    object
61 GarageCars      1460 non-null    int64
62 GarageArea      1460 non-null    int64
63 GarageQual      1379 non-null    object
64 GarageCond      1379 non-null    object
65 PavedDrive     1460 non-null    object
66 WoodDeckSF     1460 non-null    int64
67 OpenPorchSF    1460 non-null    int64
68 EnclosedPorch   1460 non-null    int64
69 3SsnPorch       1460 non-null    int64
70 ScreenPorch     1460 non-null    int64
71 PoolArea        1460 non-null    int64
72 PoolQC          7 non-null      object
73 Fence            281 non-null    object
74 MiscFeature     54 non-null     object
75 MiscVal          1460 non-null    int64
76 MoSold           1460 non-null    int64
77 YrSold           1460 non-null    int64
78 SaleType         1460 non-null    object
79 SaleCondition    1460 non-null    object
80 SalePrice        1460 non-null    int64
dtypes: float64(3), int64(35), object(43)
memory usage: 924.0+ KB

```

----- TEST DATA INFO -----

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1459 entries, 0 to 1458
Data columns (total 80 columns):
 #   Column          Non-Null Count  Dtype  
--- 
 0   Id              1459 non-null    int64  
 1   MSSubClass       1459 non-null    int64  
 2   MSZoning        1455 non-null    object  
 3   LotFrontage     1232 non-null    float64 
 4   LotArea          1459 non-null    int64  
 5   Street           1459 non-null    object  
 6   Alley            107 non-null     object  
 7   LotShape          1459 non-null    object  
 8   LandContour      1459 non-null    object  
 9   Utilities         1457 non-null    object  
 10  LotConfig         1459 non-null    object  
 11  LandSlope         1459 non-null    object  
 12  Neighborhood      1459 non-null    object  
 13  Condition1        1459 non-null    object  
 14  Condition2        1459 non-null    object  
 15  BldgType          1459 non-null    object  
 16  HouseStyle         1459 non-null    object  
 17  OverallQual       1459 non-null    int64  
 18  OverallCond       1459 non-null    int64  
 19  YearBuilt          1459 non-null    int64  
 20  YearRemodAdd      1459 non-null    int64  
 21  RoofStyle          1459 non-null    object  
 22  RoofMatl           1459 non-null    object  
 23  Exterior1st        1458 non-null    object 

```

```
24  Exterior2nd    1458 non-null  object
25  MasVnrType     565 non-null   object
26  MasVnrArea     1444 non-null   float64
27  ExterQual      1459 non-null   object
28  ExterCond      1459 non-null   object
29  Foundation     1459 non-null   object
30  BsmtQual       1415 non-null   object
31  BsmtCond       1414 non-null   object
32  BsmtExposure   1415 non-null   object
33  BsmtFinType1   1417 non-null   object
34  BsmtFinSF1     1458 non-null   float64
35  BsmtFinType2   1417 non-null   object
36  BsmtFinSF2     1458 non-null   float64
37  BsmtUnfSF      1458 non-null   float64
38  TotalBsmtSF    1458 non-null   float64
39  Heating         1459 non-null   object
40  HeatingQC       1459 non-null   object
41  CentralAir      1459 non-null   object
42  Electrical      1459 non-null   object
43  1stFlrSF        1459 non-null   int64
44  2ndFlrSF        1459 non-null   int64
45  LowQualFinSF   1459 non-null   int64
46  GrLivArea       1459 non-null   int64
47  BsmtFullBath   1457 non-null   float64
48  BsmtHalfBath   1457 non-null   float64
49  FullBath        1459 non-null   int64
50  HalfBath        1459 non-null   int64
51  BedroomAbvGr   1459 non-null   int64
52  KitchenAbvGr   1459 non-null   int64
53  KitchenQual     1458 non-null   object
54  TotRmsAbvGrd   1459 non-null   int64
55  Functional      1457 non-null   object
56  Fireplaces       1459 non-null   int64
57  FireplaceQu     729 non-null   object
58  GarageType       1383 non-null   object
59  GarageYrBlt     1381 non-null   float64
60  GarageFinish     1381 non-null   object
61  GarageCars       1458 non-null   float64
62  GarageArea       1458 non-null   float64
63  GarageQual       1381 non-null   object
64  GarageCond       1381 non-null   object
65  PavedDrive       1459 non-null   object
66  WoodDeckSF       1459 non-null   int64
67  OpenPorchSF      1459 non-null   int64
68  EnclosedPorch    1459 non-null   int64
69  3SsnPorch        1459 non-null   int64
70  ScreenPorch      1459 non-null   int64
71  PoolArea         1459 non-null   int64
72  PoolQC           3 non-null    object
73  Fence             290 non-null   object
74  MiscFeature      51 non-null    object
75  MiscVal          1459 non-null   int64
76  MoSold            1459 non-null   int64
77  YrSold            1459 non-null   int64
78  SaleType          1458 non-null   object
79  SaleCondition     1459 non-null   object
dtypes: float64(11), int64(26), object(43)
memory usage: 912.0+ KB
```

```
In [38]: # Count missing values per column
missing = train.isnull().sum()
missing = missing[missing > 0].sort_values(ascending=False)

print(f"Total columns with missing values: {len(missing)}")
missing.head(20)
```

Total columns with missing values: 19

```
Out[38]: PoolQC      1453
MiscFeature    1406
Alley        1369
Fence        1179
MasVnrType     872
FireplaceQu   690
LotFrontage     259
GarageType      81
GarageYrBlt     81
GarageFinish     81
GarageQual      81
GarageCond      81
BsmtFinType2    38
BsmtExposure    38
BsmtFinType1    37
BsmtCond        37
BsmtQual        37
MasVnrArea        8
Electrical        1
dtype: int64
```

We will summarize numerical statistics using `.describe()` to check data spread, means, and possible outliers.

```
In [39]: train.describe().T.head(15)
```

Out [39] :

		count	mean	std	min	25%	50%
	Id	1460.000	730.500	421.610	1.000	365.750	730.500
	MSSubClass	1460.000	56.897	42.301	20.000	20.000	50.000
	LotFrontage	1201.000	70.050	24.285	21.000	59.000	69.000
	LotArea	1460.000	10516.828	9981.265	1300.000	7553.500	9478.500
	OverallQual	1460.000	6.099	1.383	1.000	5.000	6.000
	OverallCond	1460.000	5.575	1.113	1.000	5.000	5.000
	YearBuilt	1460.000	1971.268	30.203	1872.000	1954.000	1973.000
	YearRemodAdd	1460.000	1984.866	20.645	1950.000	1967.000	1994.000
	MasVnrArea	1452.000	103.685	181.066	0.000	0.000	0.000
	BsmtFinSF1	1460.000	443.640	456.098	0.000	0.000	383.500
	BsmtFinSF2	1460.000	46.549	161.319	0.000	0.000	0.000
	BsmtUnfSF	1460.000	567.240	441.867	0.000	223.000	477.500
	TotalBsmtSF	1460.000	1057.429	438.705	0.000	795.750	991.500
	1stFlrSF	1460.000	1162.627	386.588	334.000	882.000	1087.000
	2ndFlrSF	1460.000	346.992	436.528	0.000	0.000	0.000

Understanding the Target Variable (Saleprice) :

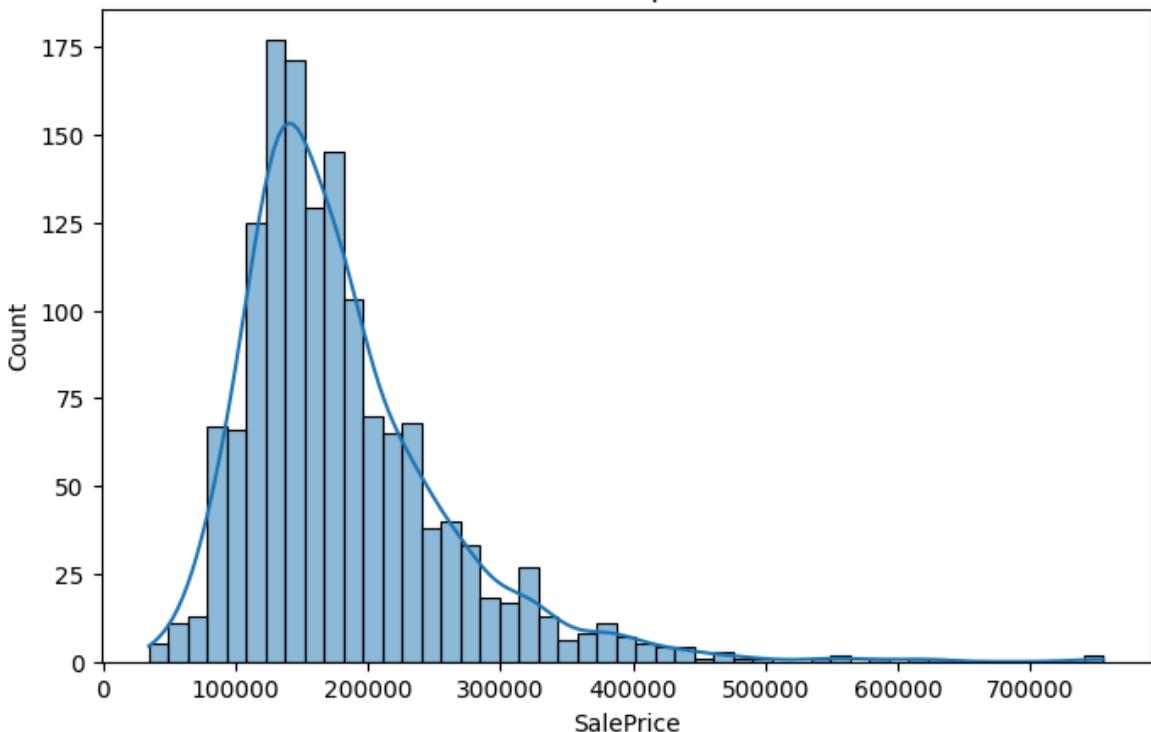
Insights:

- The price distribution is right-skewed — most houses are moderately priced with a few expensive outliers.
- 75% of all houses cost below ~214,000 USD.
- Mean > Median confirms a right-skewed distribution.

In [40] :

```
plt.figure(figsize=(8,5))
sns.histplot(train['SalePrice'], kde=True)
plt.title('Distribution of saleprice of Houses')
plt.show()
print(train['SalePrice'].describe())
```

Distribution of saleprice of Houses



```

count      1460.000
mean      180921.196
std       79442.503
min      34900.000
25%     129975.000
50%     163000.000
75%     214000.000
max      755000.000
Name: SalePrice, dtype: float64

```

```
In [41]: corr = train.select_dtypes(include=[np.number]).corr() ['SalePrice'].sort_
corr.head(15)
```

```

Out[41]: SalePrice      1.000
OverallQual      0.791
GrLivArea        0.709
GarageCars        0.640
GarageArea        0.623
TotalBsmtSF      0.614
1stFlrSF         0.606
FullBath          0.561
TotRmsAbvGrd     0.534
YearBuilt         0.523
YearRemodAdd      0.507
GarageYrBlt       0.486
MasVnrArea        0.477
Fireplaces        0.467
BsmtFinSF1        0.386
Name: SalePrice, dtype: float64

```

The above step also helps us realize why a **log transformation** will later improve model stability and reduce skewness.

Visual Exploratory Data Analysis (EDA) for (Numeric features)

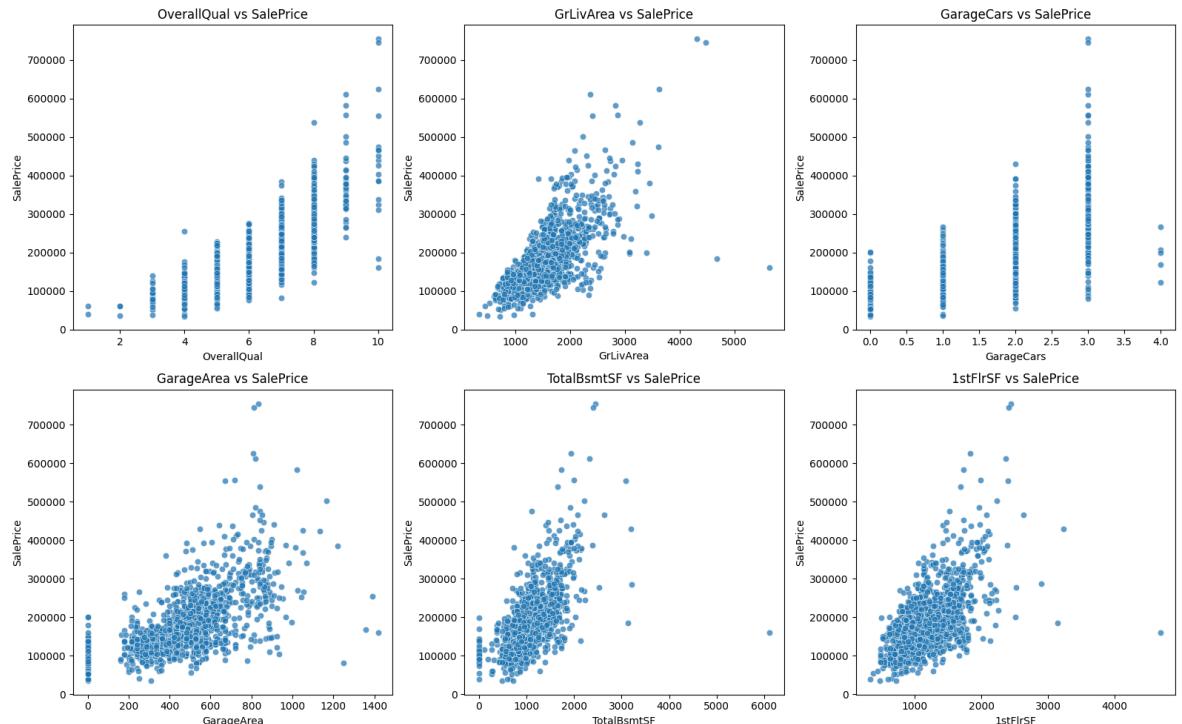
We compute correlations between numeric features and SalePrice. Top correlated features:

- OverallQual
- GrLivArea
- GarageCars
- TotalBsmtSF
- 1stFlrSF

Let's visualize how the top correlated numeric features relate to **SalePrice**.

```
In [42]: top_features = ['OverallQual', 'GrLivArea', 'GarageCars', 'GarageArea', 'TotalBsmtSF', '1stFlrSF']

plt.figure(figsize=(16, 10))
for i, feature in enumerate(top_features[:6]):
    plt.subplot(2, 3, i + 1)
    sns.scatterplot(data=train, x=feature, y='SalePrice', alpha=0.7)
    plt.title(f'{feature} vs SalePrice')
plt.tight_layout()
plt.show()
```



Above, we can already see that:

- Higher **OverallQual** and **GrLivArea** strongly increase SalePrice.
- There are potential **outliers** — very large houses sold at lower prices.

Outlier Handling & Categorical EDA

Removing outliers:

```
In [43]: # Identify potential outliers
outliers = train[(train['GrLivArea'] > 4000) & (train['SalePrice'] < 30000)]
display(outliers[['Id', 'GrLivArea', 'SalePrice']])

# Remove them from the training set
train = train.drop(outliers.index)

print(f"New train shape after removing outliers: {train.shape}")
```

	Id	GrLivArea	SalePrice
	523	524	4676
	1298	1299	184750
		5642	160000

New train shape after removing outliers: (1458, 81)

We detected **extreme outliers**: houses with GrLivArea > 4000 but low SalePrice < 300000.

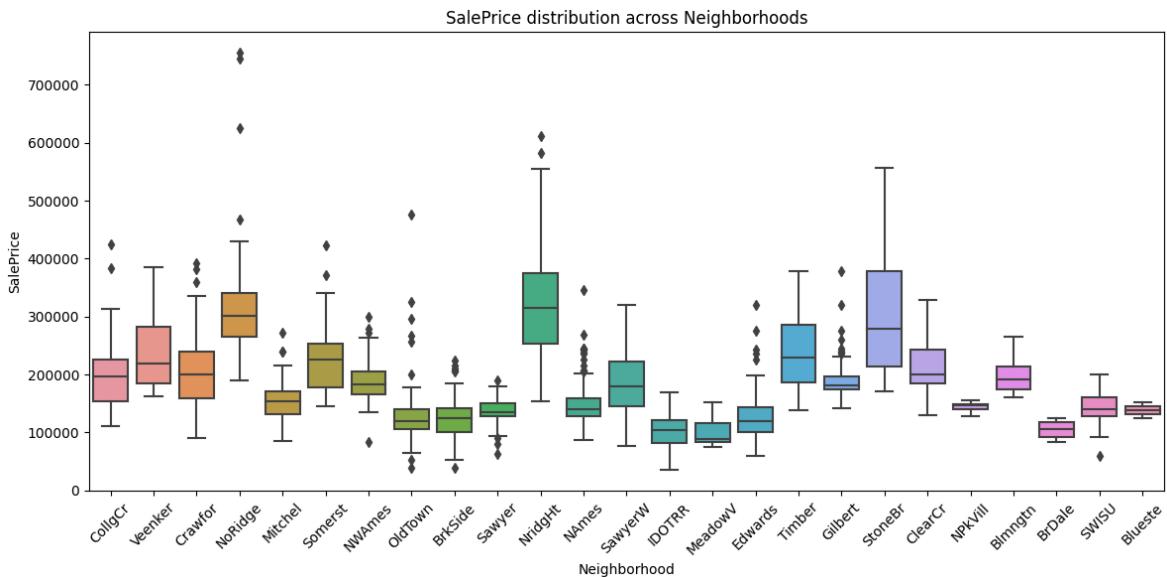
These can harm the model and introduce bias, so we remove them to improve model generalization.

Exploring categorical data:

We will visualize relationships between categorical variables and SalePrice using

Boxplots:

```
In [44]: # Plot settings for exploring Categorical features
plt.figure(figsize=(14,6))
sns.boxplot(x='Neighborhood', y='SalePrice', data=train)
plt.xticks(rotation=45)
plt.title('SalePrice distribution across Neighborhoods')
plt.show()
```



Insights:

After reviewing the above boxplot it can be seen that:

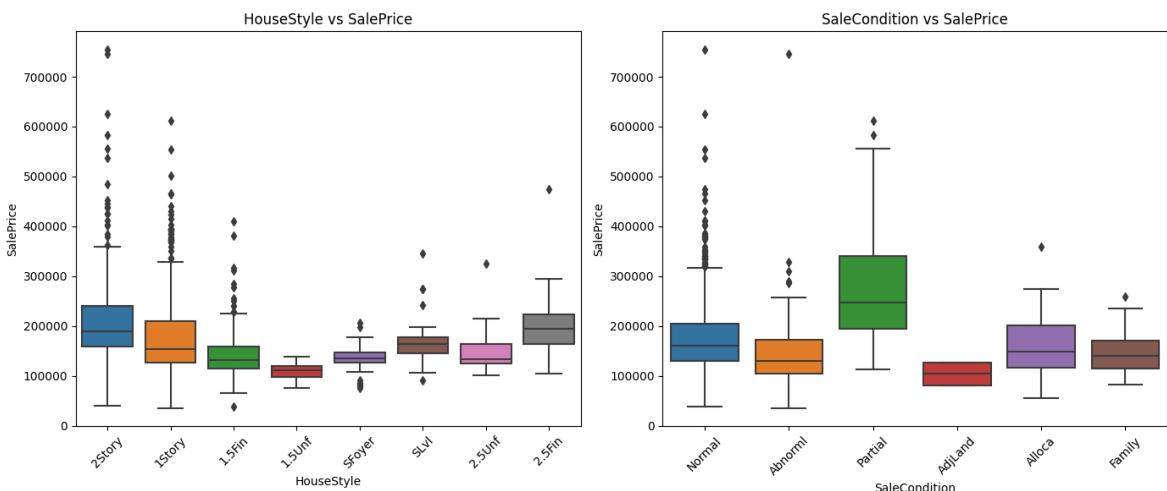
- Neighborhood**: Some neighborhoods consistently have higher prices.

```
In [45]: fig, axes = plt.subplots(1, 2, figsize=(14,6))

sns.boxplot(x='HouseStyle', y='SalePrice', data=train, ax=axes[0])
axes[0].set_title('HouseStyle vs SalePrice')
axes[0].tick_params(axis='x', rotation=45)

sns.boxplot(x='SaleCondition', y='SalePrice', data=train, ax=axes[1])
axes[1].set_title('SaleCondition vs SalePrice')
axes[1].tick_params(axis='x', rotation=45)

plt.tight_layout()
plt.show()
```



Insights 🧠

- HouseStyle**: 2Story and 1Story homes tend to have higher prices, while 1.5Fin and 1.5Unf styles are cheaper.
- SaleCondition**: 'Partial' (new construction) sales are the most expensive.

These patterns confirm **categorical features are strong predictors**.

Feature Engineering and Data Prep

Handling Missing Values and Imputation:

- Columns where NaN means "None" (e.g., Alley , Fence , PoolQC) are filled with " None ".
- Numeric basement/garage values are replaced with 0 .
- LotFrontage is filled using the **median** per Neighborhood.
- GarageYrBlt missing → 0 .
- Electrical missing → value from **mode**.

```
In [46]: import numpy as np
import pandas as pd

train = pd.read_csv('/kaggle/input/house-prices-advanced-regression-techni
test = pd.read_csv('/kaggle/input/house-prices-advanced-regression-techni

# Workin on a copy
df_train = train.copy()

# 1) Columns where NaN means "None"
none_cols = [
    'Alley', 'BsmtQual', 'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFin
    'FireplaceQu', 'GarageType', 'GarageFinish', 'GarageQual', 'GarageCond',
    'PoolQC', 'Fence', 'MiscFeature', 'MasVnrType'
]
for c in none_cols:
    if c in df_train.columns:
        df_train[c] = df_train[c].fillna('None')

# Basement/garage
zero_cols = ['BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF',
            'BsmtFullBath', 'BsmtHalfBath', 'GarageCars', 'GarageArea', 'Mas
for c in zero_cols:
    if c in df_train.columns:
        df_train[c] = df_train[c].fillna(0)

# LotFrontage: impute by Neighborhood median
if 'LotFrontage' in df_train.columns and 'Neighborhood' in df_train.colum
    df_train['LotFrontage'] = df_train.groupby('Neighborhood')['LotFrontage']
        lambda s: s.fillna(s.median())
    )

# GarageYrBlt: missing means no garage
if 'GarageYrBlt' in df_train.columns:
    df_train['GarageYrBlt'] = df_train['GarageYrBlt'].fillna(0)

# Electrical: single missing
```

```

if 'Electrical' in df_train.columns:
    df_train['Electrical'] = df_train['Electrical'].fillna(df_train['Elec'])

# sanity check of remaining NA
na_left = df_train.isna().sum()
na_left = na_left[na_left>0].sort_values(ascending=False)
print("Still missing after rules:")
print(na_left)

```

Still missing after rules:

Series([], dtype: int64)

Above output shows that there are **no missing values** anymore.

Log-Transforming the target variable SalePrice:

We apply log transformation to SalePrice using np.log1p() because,

- it reduces the influence of extremely high-priced houses.
- Makes data more normally distributed.
- Helps linear models handle target variability better.

```

In [47]: # Create a new target variable (log-transformed SalePrice)
y = np.log1p(train['SalePrice']) # log(1 + SalePrice)

# Drop SalePrice and Id to get features
X = train.drop(columns=['SalePrice', 'Id'])

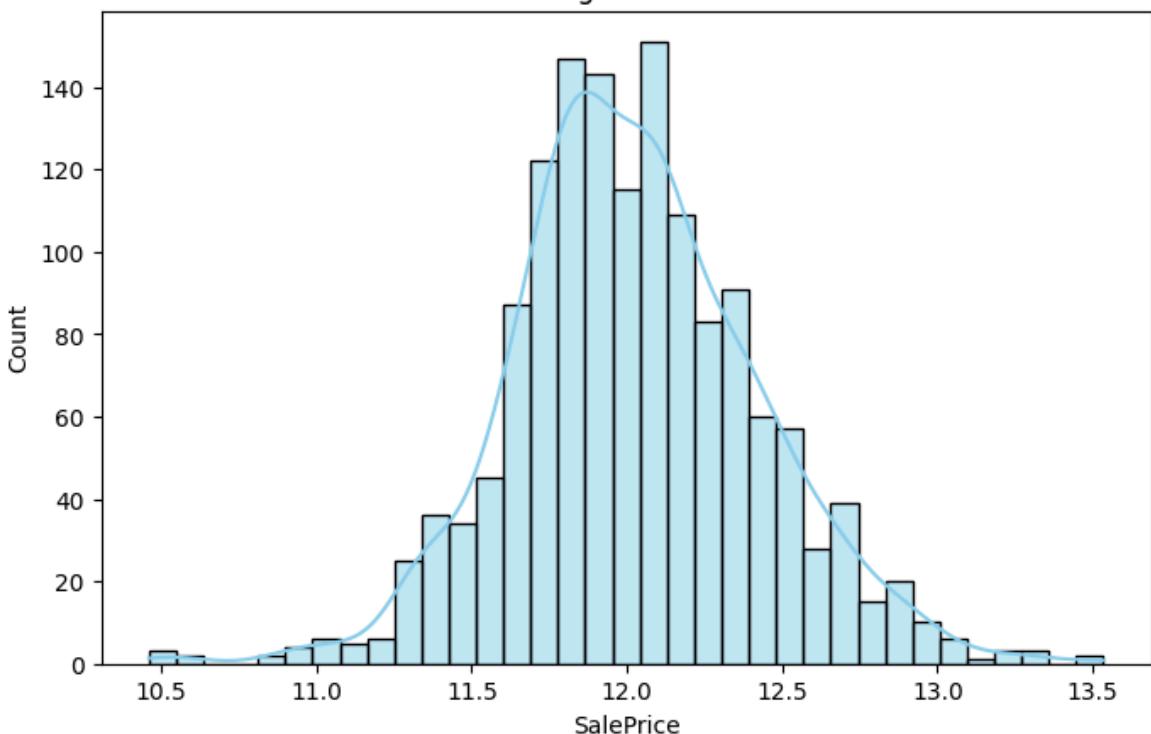
# Check transformation visually
import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(8,5))
sns.histplot(y, kde=True, color='skyblue')
plt.title('Distribution of Log-Transformed SalePrice')
plt.show()

print("y mean:", y.mean(), "y std:", y.std())

```

Distribution of Log-Transformed SalePrice



y mean: 12.024057394918406 y std: 0.3994492733225068

Now the distribution of the target `SalePrice` above is not right-skewed anymore when it is log-transformed.

Encoding the categorical values:

We will separate numeric and categorical columns:

- Apply ordinal mappings to quality-related features (e.g., ExterQual, KitchenQual, BsmtQual, etc.).
- Remaining categorical features are one-hot encoded.

This encoding allows both linear and non linear models to process mixed data types correctly.

```
In [48]: from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer

# Create a copy for safety
X = train.drop(columns=['SalePrice', 'Id']).copy()

# Ordinal mappings (keep meaningful order)
qual_map = {'None':0, 'Po':1, 'Fa':2, 'TA':3, 'Gd':4, 'Ex':5}
bsmt_exp_map = {'None':0, 'No':0, 'Mn':1, 'Av':2, 'Gd':3}
bsmt_fin_map = {'None':0, 'Unf':1, 'LwQ':2, 'Rec':3, 'BLQ':4, 'ALQ':5, 'GLQ':6}
paved_map = {'N':0, 'P':1, 'Y':2}
bin_map = {'N':0, 'Y':1}
functional_map = {'Sal':0, 'Sev':1, 'Maj2':2, 'Maj1':3, 'Mod':4, 'Min2':5, 'Min':6}
```

```

ordinal_maps = {
    'ExterQual': qual_map, 'ExterCond': qual_map,
    'BsmtQual': qual_map, 'BsmtCond': qual_map, 'BsmtExposure': bsmt_exp_
    'BsmtFinType1': bsmt_fin_map, 'BsmtFinType2': bsmt_fin_map,
    'HeatingQC': qual_map, 'KitchenQual': qual_map,
    'FireplaceQu': qual_map, 'GarageQual': qual_map, 'GarageCond': qual_m
    'PoolQC': qual_map, 'PavedDrive': paved_map, 'CentralAir': bin_map,
    'Functional': functional_map
}

for col, mapping in ordinal_maps.items():
    if col in X.columns:
        X[col] = X[col].map(mapping).astype('float64')

# Split remaining columns by dtype
num_cols = X.select_dtypes(include=['int64', 'float64']).columns.tolist()
cat_cols = X.select_dtypes(include=['object']).columns.tolist()

print(f"Numeric columns: {len(num_cols)} | Categorical columns: {len(cat_")

# Preprocess: imputers + one-hot encoder
numeric_proc = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median'))
])

categorical_proc = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('onehot', OneHotEncoder(handle_unknown='ignore', sparse_output=False)
])

preprocess = ColumnTransformer(
    transformers=[
        ('num', numeric_proc, num_cols),
        ('cat', categorical_proc, cat_cols)
    ]
)

print("Encoding setup complete – ready for model training.")

```

Numeric columns: 52 | Categorical columns: 27
 Encoding setup complete – ready for model training.

Linear Model Training

Baseline model training using Ridge Regression:

```

In [49]: # Baseline Ridge Regression with 10-Fold Cross-Validation

from sklearn.linear_model import Ridge
from sklearn.model_selection import KFold, cross_val_score
from sklearn.pipeline import Pipeline
import numpy as np

# Create a Ridge regression model inside a pipeline
ridge = Pipeline(steps=[
    ('preprocessor', preprocess),

```

```
( 'model', Ridge(alpha=10.0, random_state=42))
]

# Define 10-fold cross-validation setup
cv = KFold(n_splits=10, shuffle=True, random_state=42)

# Evaluate using Root Mean Squared Log Error
scores = cross_val_score(
    ridge, X, y,
    scoring='neg_root_mean_squared_error',
    cv=cv,
    n_jobs=-1
)

print(f"Ridge 10-Fold CV log-RMSE: {scores.mean():.4f} ± {scores.std():.4f}")

```

Ridge 10-Fold CV log-RMSE: 0.1374 ± 0.0379

Insights:

- CV log-RMSE: 0.1374 ± 0.0379
- This means we have now achieved a **13.7% relative prediction error**

Hyperparameter tuning

Now that we have the baseline model ready, we can test multiple alpha (regularization) values using `RidgeCV` and `LassoCV`

```
In [50]: from sklearn.linear_model import RidgeCV, LassoCV

# Define range of alpha values to test
alphas = np.logspace(-3, 3, 50) # from 0.001 to 1000

# RidgeCV
ridge_cv = RidgeCV(alphas=alphas, scoring='neg_root_mean_squared_error',
ridge_cv.fit(preprocess.fit_transform(X), y)
print(f"Best Ridge alpha: {ridge_cv.alpha:.4f}")

#LassoCV for comparison
lasso_cv = LassoCV(alphas=alphas, cv=10, random_state=42, max_iter=10000)
lasso_cv.fit(preprocess.fit_transform(X), y)
print(f"Best Lasso alpha: {lasso_cv.alpha:.4f}")
```

Best Ridge alpha: 10.9854

Best Lasso alpha: 0.0010

Insights:

- Ridge $\alpha \approx 10.9854$
- Lasso $\alpha \approx 0.0010$

Also note that the std deviation is still 0.0379 , we will also try to improve this.

Rechecking Cross-Validation:

```
In [51]: from sklearn.model_selection import KFold, cross_val_score
from sklearn.linear_model import Ridge, Lasso
from sklearn.pipeline import Pipeline
import numpy as np

best_alpha_ridge = 10.9854
best_alpha_lasso = 0.0010

ridge_best = Pipeline([
    ('preprocessor', preprocess),
    ('model', Ridge(alpha=best_alpha_ridge, random_state=42))
])

lasso_best = Pipeline([
    ('preprocessor', preprocess),
    ('model', Lasso(alpha=best_alpha_lasso, random_state=42, max_iter=200)
])

cv = KFold(n_splits=10, shuffle=True, random_state=42)

ridge_scores = cross_val_score(ridge_best, X, y,
                               scoring='neg_root_mean_squared_error',
                               cv=cv, n_jobs=-1)
lasso_scores = cross_val_score(lasso_best, X, y,
                               scoring='neg_root_mean_squared_error',
                               cv=cv, n_jobs=-1)

print(f"Ridge (\alpha={best_alpha_ridge:.4f}) CV: {-ridge_scores.mean():.4f}")
print(f"Lasso (\alpha={best_alpha_lasso:.4f}) CV: {-lasso_scores.mean():.4f}")

Ridge (\alpha=10.9854) CV: 0.1374 ± 0.0380
Lasso (\alpha=0.0010) CV: 0.1375 ± 0.0426
```

Insights:

- Ridge CV log-RMSE: 0.1374
- Lasso CV log-RMSE: 0.1375

Hence, it can be noted that both perform similarly when checked with cross validation.

Fine-tuning lasso again

- to check if the true best alpha(α) for Lasso is even smaller (<0.001)

For this, we will test with smaller alpha values ranging from 0.00001 to 0.01

```
In [52]: from sklearn.linear_model import LassoCV
import numpy as np

# Transform once for speed
X_mm = preprocess.fit_transform(X)

alphas_fine = np.logspace(-5, -2, 30) # 1e-5 ... 1e-2
lasso_cv_fine = LassoCV(alphas=alphas_fine, cv=10, random_state=42, max_i
lasso_cv_fine.fit(X_mm, y)
```

```
best_alpha_lasso_refined = float(lasso_cv_fine.alpha_)
print("Refined best Lasso alpha:", best_alpha_lasso_refined)
```

Refined best Lasso alpha: 0.0005736152510448681

Because the refined Lasso printed an alpha <0.001, we will re-evaluate Lasso with the new alpha 0.0005736152510448681

```
In [53]: from sklearn.model_selection import KFold, cross_val_score
from sklearn.linear_model import Lasso
from sklearn.pipeline import Pipeline

lasso_best = Pipeline([
    ('preprocessor', preprocess),
    ('model', Lasso(alpha=best_alpha_lasso_refined, random_state=42, max_iter=10))
])

cv = KFold(n_splits=10, shuffle=True, random_state=42)
lasso_scores = cross_val_score(lasso_best, X, y,
                               scoring='neg_root_mean_squared_error',
                               cv=cv, n_jobs=-1)
print(f'Lasso (α={best_alpha_lasso_refined:.6f}) CV: {-lasso_scores.mean():.4f}')
Lasso (α=0.000574) CV: 0.1342 ± 0.0439
```

- So with the updated alpha α=0.000574 that is far lesser than the older one, it confirms that a smaller α helped the Lasso model to learn more detail without overfitting
- We also improved the relative error in prediction to 0.1342 (approx. 13.4%) which was earlier 0.1375

Hence, we can say that this is our best performing linear model.

Training and Evaluating Non Linear models

Now, We will train the non linear models LightGBM and XGBoost with early stopping and CV and then generate OOF (out-of-fold) predictions for both models.

```
In [54]: import numpy as np
from sklearn.model_selection import KFold
from sklearn.metrics import mean_squared_error

# Reuse preprocessed features
X_mm = preprocess.fit_transform(X)

def rmse_log(y_true, y_pred):
    return np.sqrt(np.mean((y_true - y_pred)**2))

kf = KFold(n_splits=10, shuffle=True, random_state=42)

# ----- LightGBM -----
import lightgbm as lgb
lgb_params = dict(
    n_estimators=5000,
```

```

learning_rate=0.01,
num_leaves=31,
subsample=0.8,
colsample_bytree=0.8,
reg_alpha=0.0,
reg_lambda=0.0,
random_state=42
)

oof_lgb = np.zeros(len(X))
lgb_best_iters = []

for tr, va in kf.split(X_mm):
    X_tr, X_va = X_mm[tr], X_mm(va)
    y_tr, y_va = y.iloc[tr], y.iloc(va]

    model_lgb = lgb.LGBMRegressor(**lgb_params)
    model_lgb.fit(
        X_tr, y_tr,
        eval_set=[(X_va, y_va)],
        eval_metric='rmse',
        callbacks=[lgb.early_stopping(stopping_rounds=200, verbose=False)
    )
    oof_lgb[va] = model_lgb.predict(X_va, num_iteration=model_lgb.best_it
    lgb_best_iters.append(model_lgb.best_iteration_)

rmse_lgb = rmse_log(y, oof_lgb)
print(f"LightGBM 10-fold OOF log-RMSE: {rmse_lgb:.4f} (avg best_iter ≈ {i

# ----- XGBoost -----
from xgboost import XGBRegressor
xgb_params = dict(
    n_estimators=6000,
    learning_rate=0.01,
    max_depth=6,
    subsample=0.8,
    colsample_bytree=0.8,
    reg_alpha=0.0,
    reg_lambda=1.0,
    objective='reg:squarederror',
    random_state=42,
    n_jobs=-1
)

oof_xgb = np.zeros(len(X))
xgb_best_iters = []

for tr, va in kf.split(X_mm):
    X_tr, X_va = X_mm[tr], X_mm(va)
    y_tr, y_va = y.iloc[tr], y.iloc(va)

    model_xgb = XGBRegressor(**xgb_params)
    model_xgb.fit(
        X_tr, y_tr,
        eval_set=[(X_va, y_va)],
        eval_metric='rmse',
        verbose=False,
        early_stopping_rounds=200
    )
    best_iter = model_xgb.best_iteration if model_xgb.best_iteration is n

```

```
oof_xgb[va] = model_xgb.predict(X_va, iteration_range=(0, best_iter))
xgb_best_iters.append(best_iter)

rmse_xgb = rmse_log(y, oof_xgb)
print(f"XGBoost 10-fold OOF log-RMSE: {rmse_xgb:.4f} (avg best_iter ≈ {in
```

```
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.001051 seconds.  
You can set `force_row_wise=true` to remove the overhead.  
And if memory is not enough, you can set `force_col_wise=true`.  
[LightGBM] [Info] Total Bins 3341  
[LightGBM] [Info] Number of data points in the train set: 1314, number of used features: 147  
[LightGBM] [Info] Start training from score 12.025324  
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000929 seconds.  
You can set `force_row_wise=true` to remove the overhead.  
And if memory is not enough, you can set `force_col_wise=true`.  
[LightGBM] [Info] Total Bins 3327  
[LightGBM] [Info] Number of data points in the train set: 1314, number of used features: 149  
[LightGBM] [Info] Start training from score 12.028659  
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.001049 seconds.  
You can set `force_row_wise=true` to remove the overhead.  
And if memory is not enough, you can set `force_col_wise=true`.  
[LightGBM] [Info] Total Bins 3333  
[LightGBM] [Info] Number of data points in the train set: 1314, number of used features: 147  
[LightGBM] [Info] Start training from score 12.021956  
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.001186 seconds.  
You can set `force_row_wise=true` to remove the overhead.  
And if memory is not enough, you can set `force_col_wise=true`.  
[LightGBM] [Info] Total Bins 3332  
[LightGBM] [Info] Number of data points in the train set: 1314, number of used features: 148  
[LightGBM] [Info] Start training from score 12.019795  
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000850 seconds.  
You can set `force_row_wise=true` to remove the overhead.  
And if memory is not enough, you can set `force_col_wise=true`.  
[LightGBM] [Info] Total Bins 3317  
[LightGBM] [Info] Number of data points in the train set: 1314, number of used features: 145  
[LightGBM] [Info] Start training from score 12.020663  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.001309 seconds.  
You can set `force_col_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 3320  
[LightGBM] [Info] Number of data points in the train set: 1314, number of used features: 144  
[LightGBM] [Info] Start training from score 12.026297  
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.001081 seconds.  
You can set `force_row_wise=true` to remove the overhead.  
And if memory is not enough, you can set `force_col_wise=true`.  
[LightGBM] [Info] Total Bins 3332  
[LightGBM] [Info] Number of data points in the train set: 1314, number of used features: 147  
[LightGBM] [Info] Start training from score 12.029436  
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.001147 seconds.  
You can set `force_row_wise=true` to remove the overhead.  
And if memory is not enough, you can set `force_col_wise=true`.  
[LightGBM] [Info] Total Bins 3329
```

```
[LightGBM] [Info] Number of data points in the train set: 1314, number of used features: 147
[LightGBM] [Info] Start training from score 12.022123
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.001099 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 3331
[LightGBM] [Info] Number of data points in the train set: 1314, number of used features: 147
[LightGBM] [Info] Start training from score 12.023877
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000802 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 3335
[LightGBM] [Info] Number of data points in the train set: 1314, number of used features: 146
[LightGBM] [Info] Start training from score 12.022444
LightGBM 10-fold OOF log-RMSE: 0.1247 (avg best_iter ≈ 1879)
XGBoost 10-fold OOF log-RMSE: 0.1241 (avg best_iter ≈ 2251)
```

Insights:

Now we have the OOF predictions for **LightGBM and XGBoost** which is,

- log RMSE: 0.1247 **approximately 12.47%** for LightGBM
- and log RMSE: 0.1241 **approximately 12.41%** for XGBoost

Next step is to combine all four of the models' OOF predictions and blend them together for getting the best possible **Out-Of-Fold (OOF) log RMSE**.

Generating OOF Predictions for Ridge and Lasso

We have to generate the OOF predictions for our Linear models because we already have the OOF log RMSE for the non linear models (LightGBM and XGBoost)

- we are doing this to match the prediction metrics of all 4 models, so that later we can ensemble them together and get the final OOF log RMSE.
- This will ensure all four models (**Ridge, Lasso, LightGBM, XGBoost**) are evaluated using consistent folds.

```
In [55]: import numpy as np
from sklearn.model_selection import KFold
from sklearn.linear_model import Ridge, Lasso
from sklearn.pipeline import Pipeline

# 10-Fold Cross Validation setup (same folds for all models for fairness)
kf = KFold(n_splits=10, shuffle=True, random_state=42)

# Create Ridge and Lasso pipelines with tuned alphas
ridge_best = Pipeline([
    ('preprocessor', preprocess),
```

```
( 'model', Ridge(alpha=10.9854, random_state=42))
])

lasso_best = Pipeline([
    ('preprocessor', preprocess),
    ('model', Lasso(alpha=0.0005736152510448681, random_state=42, max_ite
])

# Empty arrays to store out-of-fold predictions (same length as training
oof_ridge = np.zeros(len(X))
oof_lasso = np.zeros(len(X))

# Generate OOF predictions for Ridge and Lasso
for tr, va in kf.split(X):
    ridge_best.fit(X.iloc[tr], y.iloc[tr])
    oof_ridge[va] = ridge_best.predict(X.iloc[va])

    lasso_best.fit(X.iloc[tr], y.iloc[tr])
    oof_lasso[va] = lasso_best.predict(X.iloc[va])

# RMSE in log space (≈ RMSLE on original scale)
def rmse_log(y_true, y_pred):
    return np.sqrt(np.mean((y_true - y_pred)**2))

ridge_rmse = rmse_log(y, oof_ridge)
lasso_rmse = rmse_log(y, oof_lasso)
print("Ridge OOF log RMSE", ridge_rmse),
print("Lasso OOF log RMSE", lasso_rmse)
```

Ridge OOF log RMSE 0.14259419340182172
Lasso OOF log RMSE 0.14114880217613698

Now that we have the values of **Ridge and Lasso** OOF log RMSE above at approximately 14.25% Ridge and 14.11% Lasso,

- We can now combine all four models and blend them to see if the final log RMSE delivers a lower score than all of the models individually.

Blending All Models

- we will test the best set of weights for the four models by assigning multiple grids of weights for each model.
- we will only choose the set of weights which proves to deliver the **lowest RMSLE** of all.

Note that we will be **using the OOF prediction values for all 4 models**

```
In [56]: # Candidate weights for Ridge, Lasso, LightGBM, XGBoost
candidates = [
    (0.20, 0.20, 0.30, 0.30),
    (0.15, 0.15, 0.35, 0.35),
    (0.10, 0.10, 0.40, 0.40),
    (0.25, 0.25, 0.25, 0.25),
    (0.33, 0.33, 0.17, 0.17),
    (0.10, 0.20, 0.35, 0.35),
    (0.20, 0.10, 0.35, 0.35),
```

```
[1]
best_w = None
best_score = 999

# Test each weight combination and pick the best one (lowest RMSLE)
for w in candidates:
    w_r, w_l, w_lb, w_x = w
    blend = w_r*oof_ridge + w_l*oof_lasso + w_lb*oof_lgb + w_x*oof_xgb
    score = rmse_log(y, blend)
    print(f"Weights {w} → OOF log-RMSE: {score:.4f}")
    if score < best_score:
        best_score, best_w = score, w

print("\n Best OOF blend: weights={best_w}, score={best_score:.4f}")
```

Weights (0.2, 0.2, 0.3, 0.3) → OOF log-RMSE: 0.1215
 Weights (0.15, 0.15, 0.35, 0.35) → OOF log-RMSE: 0.1207
 Weights (0.1, 0.1, 0.4, 0.4) → OOF log-RMSE: 0.1206
 Weights (0.25, 0.25, 0.25, 0.25) → OOF log-RMSE: 0.1231
 Weights (0.33, 0.33, 0.17, 0.17) → OOF log-RMSE: 0.1272
 Weights (0.1, 0.2, 0.35, 0.35) → OOF log-RMSE: 0.1206
 Weights (0.2, 0.1, 0.35, 0.35) → OOF log-RMSE: 0.1208

Best OOF blend: weights=(0.1, 0.1, 0.4, 0.4), score=0.1206

Now after reviewing the final log RMSE scores with respect to each grid of weights assigned after combining all 4 models,

- It can be said that **0.1206** is the best score until now.
- It should also be noted that **0.1206 (approx.12%)** is a better score than all of the previously predicted OOF log RMSE of each model individually.

We can now move ahead and train our best grid of weights, **weights=(0.1, 0.1, 0.4, 0.4)** on full data.

Training and Testing on entire data

Handling a Preprocessing Error Before Final Training:

While preparing the final models for full training and test prediction, I encountered a **ValueError** related to missing value imputation:

```
ValueError: Cannot use median strategy with non-  
numeric data: could not convert string to float: 'TA'
```

This occurred because some **ordinal categorical features** (like **ExterQual**, **BsmtQual**, **KitchenQual**, etc.) contained string values ('TA', 'Gd', 'Ex', etc.) in the **test set**, which conflicted with the **numeric median imputer** in the preprocessing pipeline.

Step A: Ordinal Mapping & Combined Preprocessing

In [57]: # MAP ORDINALS & FIT PREPROCESSOR ON TRAIN + TEST

```

import lightgbm as lgb
from xgboost import XGBRegressor
from sklearn.linear_model import Ridge, Lasso

# Use average best iterations from early stopping
avg_lgb_iter = int(np.mean(lgb_best_iters))
avg_xgb_iter = int(np.mean(xgb_best_iters))

# making copies
X_train_fixed = X.copy()
X_test_fixed = test.drop(columns=['Id']).copy()

# Define the same ordinal maps used earlier
qual_map      = {'Po':1, 'Fa':2, 'TA':3, 'Gd':4, 'Ex':5}
bsmt_exp_map  = {'None':0, 'No':1, 'Mn':2, 'Av':3, 'Gd':4}
bsmt_fin_map  = {'None':0, 'Unf':1, 'LwQ':2, 'Rec':3, 'BLQ':4, 'ALQ':5,
paved_map     = {'N':0, 'P':1, 'Y':2}
bin_map        = {'N':0, 'Y':1}
functional_map = {'Sal':1, 'Sev':2, 'Maj2':3, 'Maj1':4, 'Mod':5, 'Min2':6, 'Min1':7}

ordinal_maps = {
    'ExterQual': qual_map, 'ExterCond': qual_map,
    'BsmtQual': qual_map, 'BsmtCond': qual_map, 'BsmtExposure': bsmt_exp_map,
    'BsmtFinType1': bsmt_fin_map, 'BsmtFinType2': bsmt_fin_map,
    'HeatingQC': qual_map, 'KitchenQual': qual_map,
    'FireplaceQu': qual_map, 'GarageQual': qual_map, 'GarageCond': qual_map,
    'PoolQC': qual_map, 'PavedDrive': paved_map, 'CentralAir': bin_map,
    'Functional': functional_map
}

def apply_ordinal_maps_inplace(df):
    for col, mp in ordinal_maps.items():
        if col in df.columns:
            if df[col].dtype == 'O' or df[col].dtype.name == 'category':
                df[col] = df[col].fillna('None').map(mp)
            df[col] = df[col].astype('float', errors='ignore')

apply_ordinal_maps_inplace(X_train_fixed)
apply_ordinal_maps_inplace(X_test_fixed)

# Fit existing preprocessor on TRAIN + TEST combined
combined = pd.concat([X_train_fixed, X_test_fixed], axis=0)
preprocess_full = preprocess.fit(combined)

# Transform train and test with fitted preprocessor
X_full    = preprocess_full.transform(X_train_fixed)
X_test_mm = preprocess_full.transform(X_test_fixed)

print("✅ Preprocessing completed. Shapes:", X_full.shape, X_test_mm.shape)

```

✅ Preprocessing completed. Shapes: (1460, 231) (1459, 231)

Insights:

The fix for above error:

- Mapped ordinal strings (like `TA`, `Gd`, `Ex`) to numbers in both train and test sets.
- Refit the preprocessor on combined train + test data to learn all category levels.
- Transformed both datasets again for clean modeling.

Step B: Train all models on full data

We train all four models on the full dataset using tuned parameters:

- **Ridge** ($\alpha = 10.9854$)
- **Lasso** ($\alpha = 0.000574$)
- **LightGBM** (`best_iter` ≈ 1879)
- **XGBoost** (`best_iter` ≈ 2251)

In [58]:

```
# =====
# 🌱 PART 3B - TRAIN FULL MODELS (LIGHTGBM, XGBOOST, RIDGE, LASSO)
# =====

# Train LightGBM on full data
lgb_params = dict(
    n_estimators=avg_lgb_iter,
    learning_rate=0.01,
    num_leaves=31,
    subsample=0.8,
    colsample_bytree=0.8,
    reg_alpha=0.0,
    reg_lambda=0.0,
    random_state=42
)
lgb_full = lgb.LGBMRegressor(**lgb_params)
lgb_full.fit(X_full, y)

# Train XGBoost on full data
xgb_params = dict(
    n_estimators=avg_xgb_iter,
    learning_rate=0.01,
    max_depth=6,
    subsample=0.8,
    colsample_bytree=0.8,
    reg_alpha=0.0,
    reg_lambda=1.0,
    objective='reg:squarederror',
    random_state=42,
    n_jobs=-1
)
xgb_full = XGBRegressor(**xgb_params)
xgb_full.fit(X_full, y, verbose=False)

# Train Ridge and Lasso on full data
ridge_full = Ridge(alpha=10.9854, random_state=42)
ridge_full.fit(X_full, y)
```

```

lasso_full = Lasso(alpha=0.0005736152510448681, random_state=42, max_iter
lasso_full.fit(X_full, y)

print("Models trained on full data (Ridge, Lasso, LightGBM, XGBoost).")

```

[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.001229 seconds.
 You can set `force_row_wise=true` to remove the overhead.
 And if memory is not enough, you can set `force_col_wise=true`.
 [LightGBM] [Info] Total Bins 3432
 [LightGBM] [Info] Number of data points in the train set: 1460, number of used features: 150
 [LightGBM] [Info] Start training from score 12.024057
 Models trained on full data (Ridge, Lasso, LightGBM, XGBoost).

All models are now blended and trained together successfully.

Step C: Predicting Test set, Blending final Results and Creating submission file.

In this final step, We will generate predictions for the test set:

- Get log predictions from each model.
- Blend them using the best weights (0.1, 0.1, 0.4, 0.4) .
- Convert log predictions back to actual Dollar prices.

```

In [59]: # Get log predictions from each model
pred_ridge_log = ridge_full.predict(X_test_mm)
pred_lasso_log = lasso_full.predict(X_test_mm)
pred_lgb_log   = lgb_full.predict(X_test_mm)
pred_xgb_log   = xgb_full.predict(X_test_mm)

# Blend predictions using best weights
w_r, w_l, w_lb, w_x = best_w # Example: (0.1, 0.1, 0.4, 0.4)
pred_log = (w_r*pred_ridge_log +
            w_l*pred_lasso_log +
            w_lb*pred_lgb_log +
            w_x*pred_xgb_log)

# Convert log predictions back to actual dollar prices
pred = np.expm1(pred_log)
pred = np.clip(pred, 0, None)

# Create final submission file
sub = test[['Id']].copy()
sub['SalePrice'] = pred
sub.to_csv('submission.csv', index=False)

print(f"\nSaved submission.csv with weights={best_w}, OOF blend score={be

```

Saved submission.csv with weights=(0.1, 0.1, 0.4, 0.4), OOF blend score=0.1206

Above we finally received a competitive OOF blend score of all 4 models, that is:

0.1206 (approx. 12% error)

🧠 Key Learnings:

- Regularization improves linear stability (Ridge, Lasso).
- Tree-based models handle complex non-linear relationships better.
- Blending reduces overfitting and improves generalization.
- Proper preprocessing and ordinal mapping are crucial for error-free model pipelines.

UPDATE - Trying to achieve an even better relative error than 12%

Part A — Feature engineering + neighborhood target encoding + preprocess

Part A.1

```
In [60]: # FEATURE ENGINEERING (applied to BOTH train & test)

import numpy as np
import pandas as pd
from sklearn.model_selection import KFold
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.impute import SimpleImputer

# Start from your existing frames
X_base = X.copy()
test_base = test.copy()

def add_features(df):
    df = df.copy()

    # 1) Total square footage (finished areas)
    df['TotalSF'] = df.get('TotalBsmtSF', 0) + df.get('1stFlrSF', 0) + df

    # 2) Total bathrooms (full + half baths upstairs + basement)
    df['TotalBath'] =
        df.get('FullBath', 0)
        + 0.5 * df.get('HalfBath', 0)
```

```

        + df.get('BsmtFullBath', 0)
        + 0.5 * df.get('BsmtHalfBath', 0)
    )

# 3) Ages relative to sale year (how old things are)
df['Age']      = (df.get('YrSold', 0) - df.get('YearBuilt', 0)).clip
df['RemodAge'] = (df.get('YrSold', 0) - df.get('YearRemodAdd', 0)).c
df['GarageAge'] = (df.get('YrSold', 0) - df.get('GarageYrBlt', 0)).cl

# 4) Porch/Deck footprint and a binary flag
porch_cols = ['WoodDeckSF', 'OpenPorchSF', 'EnclosedPorch', '3SsnPorch',
df['PorchSF'] = df[porch_cols].sum(axis=1)
df['HasPorch'] = (df['PorchSF'] > 0).astype(int)

# 5) Interaction: quality × living area (bigger *and* better quality
df['Qual_x_GrLiv'] = df.get('OverallQual', 0) * df.get('GrLivArea', 0)

# 6) Log transforms for very skewed size features (keep originals too
for c in ['GrLivArea', 'TotalSF', 'LotArea']:
    if c in df.columns:
        df[f'log1p_{c}'] = np.log1p(df[c])

# 7) New-build hint from SaleCondition
df['IsNew'] = (df.get('SaleCondition', 'Normal') == 'Partial').astype

return df

```

```

X_fe      = add_features(X_base)
test_fe   = add_features(test_base)
print("Added engineered features. Example new cols:",
      [c for c in ['TotalSF', 'TotalBath', 'Age', 'PorchSF', 'HasPorch', 'Qual']

```

Added engineered features. Example new cols: ['TotalSF', 'TotalBath', 'Age', 'PorchSF', 'HasPorch', 'Qual']

Insights:

- `TotalSF`: because buyers pay for total finished area, not just one floor.
- `TotalBath`: more baths = higher price; half baths count as 0.5.
- `Age` / `RemodAge` / `GarageAge`: newer houses/remodels/garages usually sell higher.
- `PorchSF` / `HasPorch`: outdoor space adds value.
- `Qual×GrLiv`: large houses with high quality get a price "boost" beyond a simple sum.
- `log1p(size)`: normalizes extreme values so models learn smoother relations.
- `IsNew`: `SaleCondition='Partial'` often means new construction.

Part A.2 - Target Encoding for Neighborhood (Leak-Free)

In this step, I created a new feature called `TE_Neighborhood`, which assigns each neighborhood the average log sale price based only on training folds. I used leak-free K-Fold target encoding, meaning each row's encoded value is computed from other folds and never from itself.

```
In [61]: # LEAK-FREE TARGET ENCODING FOR Neighborhood

from sklearn.model_selection import KFold

# y is your log-transformed target from earlier steps (np.log1p(SalePrice)
kf = KFold(n_splits=10, shuffle=True, random_state=42)
global_mean = y.mean() # fallback

# Out-of-fold means for train
te_tr = pd.Series(index=X_fe.index, dtype=float)

for tr_idx, val_idx in kf.split(X_fe):
    tr_neigh = X_fe.loc[tr_idx, 'Neighborhood'].astype(str)
    val_neigh = X_fe.loc[val_idx, 'Neighborhood'].astype(str)
    means = y.iloc[tr_idx].groupby(tr_neigh).mean() # mean log(SalePrice
    te_tr.iloc[val_idx] = val_neigh.map(means).fillna(global_mean)

X_fe['TE_Neighborhood'] = te_tr

# Apply train means to test
train_means_final = y.groupby(X_fe['Neighborhood'].astype(str)).mean()
test_fe['TE_Neighborhood'] = (
    test_fe['Neighborhood'].astype(str).map(train_means_final).fillna(glo
)

print("Target-encoded Neighborhood added as TE_Neighborhood.")
```

Target-encoded Neighborhood added as `TE_Neighborhood`.

Insights:

- Encoding is done using **10-fold CV**
- Training data = out-of-fold neighborhood means
- Test data = neighborhood means computed from full train
- Global mean is used for unseen neighborhood levels

Part A.3 — Rebuilding Preprocessing After Feature Engineering

Since feature engineering added new numerical features (like `TotalSF`, `Age`, log features, TE features, etc.), I had to refit the entire preprocessing pipeline on both train + test rows combined.

```
In [62]: # APPLY ORDINAL MAPS + PREPROCESS + TRANSFORM

# Reusing ordinal quality maps
qual_map      = {'Po':1, 'Fa':2, 'TA':3, 'Gd':4, 'Ex':5}
```

```

bsmt_exp_map = {'None':0, 'No':1, 'Mn':2, 'Av':3, 'Gd':4}
bsmt_fin_map = {'None':0, 'Unf':1, 'LwQ':2, 'Rec':3, 'BLQ':4, 'ALQ':5,
paved_map = {'N':0, 'P':1, 'Y':2}
bin_map = {'N':0, 'Y':1}
functional_map = {'Sal':1, 'Sev':2, 'Maj2':3, 'Maj1':4, 'Mod':5, 'Min2':6, 'Min3':7}

ordinal_maps = {
    'ExterQual': qual_map, 'ExterCond': qual_map,
    'BsmtQual': qual_map, 'BsmtCond': qual_map, 'BsmtExposure': bsmt_exp_map,
    'BsmtFinType1': bsmt_fin_map, 'BsmtFinType2': bsmt_fin_map,
    'HeatingQC': qual_map, 'KitchenQual': qual_map,
    'FireplaceQu': qual_map, 'GarageQual': qual_map, 'GarageCond': qual_map,
    'PoolQC': qual_map, 'PavedDrive': paved_map, 'CentralAir': bin_map,
    'Functional': functional_map
}

def apply_ordinal_maps_inplace(df):
    for col, mp in ordinal_maps.items():
        if col in df.columns:
            if df[col].dtype == 'O' or df[col].dtype.name == 'category':
                df[col] = df[col].fillna('None').map(mp)
            df[col] = df[col].astype('float', errors='ignore')

apply_ordinal_maps_inplace(X_fe)
apply_ordinal_maps_inplace(test_fe)

# Build a fresh ColumnTransformer on the engineered frames
numeric_features = X_fe.select_dtypes(include=[np.number]).columns.to_list()
categorical_features = X_fe.select_dtypes(exclude=[np.number]).columns.to_list()

numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler(with_mean=False)) # safe for sparse output
])
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('onehot', OneHotEncoder(handle_unknown='ignore', sparse=True))
])

preprocess_v2 = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features),
    ],
    remainder='drop'
)

# Fit on train+test combined features so encoders see all categories
combined_v2 = pd.concat([X_fe, test_fe], axis=0)
preprocess_v2.fit(combined_v2)

# Transform to model-ready matrices
Xmm = preprocess_v2.transform(X_fe)
Xmm_test = preprocess_v2.transform(test_fe)

print("Shapes after FE + TE + preprocess:", Xmm.shape, Xmm_test.shape)

```

Shapes after FE + TE + preprocess: (1460, 244) (1459, 244)

PART B — Out-of-Fold Predictions for Base Models

Part B.1 - Training Base Models to Generate OOF Predictions

Here I trained five different base models using 10-fold out-of-fold (OOF) training:

- Ridge Regression
- Lasso Regression
- ElasticNet
- LightGBM
- XGBoost

I used the same fold splits for all models so their OOF predictions align row-by-row.

```
In [63]: from sklearn.model_selection import KFold
from sklearn.metrics import mean_squared_error
from sklearn.linear_model import Ridge, Lasso, ElasticNet
import lightgbm as lgb
from xgboost import XGBRegressor
import numpy as np
import pandas as pd

def rmsle(y_true, y_pred):
    return np.sqrt(mean_squared_error(y_true, y_pred))

kf = KFold(n_splits=10, shuffle=True, random_state=42)
n = Xmm.shape[0]

# storage
oof = {}
test_preds = {}

def fit_oof(name, model_fn, Xtr, ytr, Xte, folds=kf):
    oof_pred = np.zeros(n)
    te_pred = np.zeros(Xte.shape[0])
    best_iters = []

    for tr_idx, val_idx in folds.split(Xtr):
        X_tr, X_val = Xtr[tr_idx], Xtr[val_idx]
        y_tr, y_val = ytr.iloc[tr_idx], ytr.iloc[val_idx]

        m = model_fn()

        # LightGBM special fitting (early stopping)
        if isinstance(m, lgb.LGBMRegressor):
            m.fit(X_tr, y_tr,
                  eval_set=[(X_val, y_val)],
                  eval_metric='rmse',
```

```

        callbacks=[lgb.early_stopping(200, verbose=False)])
best_iters.append(m.best_iteration_)
oof_pred[val_idx] = m.predict(X_val, num_iteration=m.best_itera-
te_pred += m.predict(Xte, num_iteration=m.best_iteration_) /

# XGBoost special fitting
elif isinstance(m, XGBRegressor):
    m.fit(
        X_tr, y_tr,
        eval_set=[(X_val, y_val)],
        eval_metric='rmse',
        verbose=False,
        early_stopping_rounds=200
    )

# In new XGBoost versions:
best_iters.append(m.best_iteration)

# OOF predictions using best_iteration
oof_pred[val_idx] = m.predict(X_val, iteration_range=(0, m.be

# Test predictions averaged across folds
te_pred += m.predict(Xte, iteration_range=(0, m.best_iteratio

else:
    # linear models
    m.fit(X_tr, y_tr)
    oof_pred[val_idx] = m.predict(X_val)
    te_pred += m.predict(Xte) / folds.get_n_splits()

print(f'{name}: OOF log-RMSE = {rmsle(y, oof_pred):.5f}',
      f' | avg best_iter ~ {int(np.mean(best_iters))}' if best_iters

oof[name] = oof_pred
test_preds[name] = te_pred

# --- Linear models (strong tuning) ---
def ridge_model():
    return Ridge(alpha=8.0, random_state=42)

def lasso_model():
    return Lasso(alpha=5.7e-4, max_iter=60000, random_state=42)

def enet_model():
    return ElasticNet(alpha=1e-4, l1_ratio=0.2, max_iter=60000, random_st

fit_oof("Ridge", ridge_model, Xmm, y, Xmm_test)
fit_oof("Lasso", lasso_model, Xmm, y, Xmm_test)
fit_oof("ElasticNet", enet_model, Xmm, y, Xmm_test)

# --- LightGBM (stronger settings) ---
def lgbm_model():
    return lgb.LGBMRegressor(
        learning_rate=0.005,
        n_estimators=12000,
        num_leaves=31,
        min_data_in_leaf=10,
        subsample=0.8,

```

```
    colsample_bytree=0.8,
    reg_alpha=0.1,
    reg_lambda=0.5,
    random_state=42
)

# --- XGBoost (stronger settings) ---
def xgb_model():
    return XGBRegressor(
        learning_rate=0.008,
        n_estimators=15000,
        max_depth=4,
        min_child_weight=1,
        subsample=0.75,
        colsample_bytree=0.75,
        reg_alpha=0.001,
        reg_lambda=1.0,
        objective='reg:squarederror',
        random_state=42,
        n_jobs=-1
)

fit_oof("LightGBM", lgbm_model, Xmm, y, Xmm_test)
fit_oof("XGBoost", xgb_model, Xmm, y, Xmm_test)
```

```
Ridge: OOF log-RMSE = 0.12839
Lasso: OOF log-RMSE = 0.12810
ElasticNet: OOF log-RMSE = 0.13286
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will
be ignored. Current value: min_data_in_leaf=10
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will
be ignored. Current value: min_data_in_leaf=10
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of
testing was 0.001579 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 5349
[LightGBM] [Info] Number of data points in the train set: 1314, number of
used features: 181
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will
be ignored. Current value: min_data_in_leaf=10
[LightGBM] [Info] Start training from score 12.025324
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will
be ignored. Current value: min_data_in_leaf=10
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will
be ignored. Current value: min_data_in_leaf=10
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will
be ignored. Current value: min_data_in_leaf=10
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will
be ignored. Current value: min_data_in_leaf=10
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will
be ignored. Current value: min_data_in_leaf=10
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of
testing was 0.001611 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 5328
[LightGBM] [Info] Number of data points in the train set: 1314, number of
used features: 180
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will
be ignored. Current value: min_data_in_leaf=10
[LightGBM] [Info] Start training from score 12.028659
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will
be ignored. Current value: min_data_in_leaf=10
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will
be ignored. Current value: min_data_in_leaf=10
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will
be ignored. Current value: min_data_in_leaf=10
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will
be ignored. Current value: min_data_in_leaf=10
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will
be ignored. Current value: min_data_in_leaf=10
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of
testing was 0.001302 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true` .
[LightGBM] [Info] Total Bins 5346
[LightGBM] [Info] Number of data points in the train set: 1314, number of
used features: 183
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will
be ignored. Current value: min_data_in_leaf=10
[LightGBM] [Info] Start training from score 12.021956
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will
be ignored. Current value: min_data_in_leaf=10
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will
be ignored. Current value: min_data_in_leaf=10
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will
be ignored. Current value: min_data_in_leaf=10
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will
be ignored. Current value: min_data_in_leaf=10
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will
be ignored. Current value: min_data_in_leaf=10
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of
```

```
testing was 0.001473 seconds.  
You can set `force_row_wise=true` to remove the overhead.  
And if memory is not enough, you can set `force_col_wise=true`.  
[LightGBM] [Info] Total Bins 5335  
[LightGBM] [Info] Number of data points in the train set: 1314, number of  
used features: 180  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Info] Start training from score 12.019795  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of  
testing was 0.001647 seconds.  
You can set `force_col_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 5340  
[LightGBM] [Info] Number of data points in the train set: 1314, number of  
used features: 184  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Info] Start training from score 12.020663  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of  
testing was 0.001678 seconds.  
You can set `force_col_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 5330  
[LightGBM] [Info] Number of data points in the train set: 1314, number of  
used features: 180  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Info] Start training from score 12.026297  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of  
testing was 0.001889 seconds.  
You can set `force_col_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 5336  
[LightGBM] [Info] Number of data points in the train set: 1314, number of  
used features: 181  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Info] Start training from score 12.029436
```

```
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of  
testing was 0.001894 seconds.  
You can set `force_col_wise=true` to remove the overhead.  
[LightGBM] [Info] Total Bins 5327  
[LightGBM] [Info] Number of data points in the train set: 1314, number of  
used features: 179  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Info] Start training from score 12.022123  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of  
testing was 0.001278 seconds.  
You can set `force_row_wise=true` to remove the overhead.  
And if memory is not enough, you can set `force_col_wise=true`.  
[LightGBM] [Info] Total Bins 5330  
[LightGBM] [Info] Number of data points in the train set: 1314, number of  
used features: 180  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Info] Start training from score 12.023877  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of  
testing was 0.001044 seconds.  
You can set `force_row_wise=true` to remove the overhead.  
And if memory is not enough, you can set `force_col_wise=true`.  
[LightGBM] [Info] Total Bins 5345  
[LightGBM] [Info] Number of data points in the train set: 1314, number of  
used features: 183  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Info] Start training from score 12.022444  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
[LightGBM] [Warning] min_data_in_leaf is set=10, min_child_samples=20 will  
be ignored. Current value: min_data_in_leaf=10  
LightGBM: OOF log-RMSE = 0.12316 | avg best_iter ≈ 1947  
XGBoost: OOF log-RMSE = 0.11586 | avg best_iter ≈ 2369
```

Part B.2 — Building the Stacking Matrices

```
In [64]: import pandas as pd

stack_train = pd.DataFrame({name: oof[name] for name in oof})
stack_test = pd.DataFrame({name: test_preds[name] for name in test_preds})

print("Stack train shape:", stack_train.shape)
print("Stack test shape:", stack_test.shape)

stack_train.head()
```

Stack train shape: (1460, 5)
 Stack test shape: (1459, 5)

	Ridge	Lasso	ElasticNet	LightGBM	XGBoost
0	12.244	12.243	12.242	12.255	12.235
1	12.189	12.214	12.186	12.129	12.119
2	12.274	12.279	12.265	12.292	12.271
3	12.087	12.137	12.051	12.143	12.082
4	12.639	12.650	12.620	12.621	12.619

Insights:

Using the OOF predictions, I created two new matrices:

`stack_train` → shape (1460, 5) Contains OOF predictions from all 5 base models for training rows.

`stack_test` → shape (1459, 5) Contains averaged predictions from all 5 base models for test rows.

These two matrices will be the input to the stacking meta-model in the next part.

PART C — Final Stacked Meta-Model

In this final step, I trained a **Lasso meta-model** on top of the stacking matrices. The model learns how to combine Ridge, Lasso, ElasticNet, LightGBM, and XGBoost.

I used 10-fold CV again to generate meta-level OOF predictions and then used the OOf predictions test predictions across folds.

```
In [65]: from sklearn.linear_model import Lasso
from sklearn.model_selection import KFold
from sklearn.metrics import mean_squared_error
import numpy as np

def rmsle(y_true, y_pred):
    return np.sqrt(mean_squared_error(y_true, y_pred))
```

```

# 10-fold CV for meta-model
kf = KFold(n_splits=10, shuffle=True, random_state=42)

oof_meta = np.zeros(stack_train.shape[0])      # OOF predictions for meta-
test_meta = np.zeros(stack_test.shape[0])       # Test predictions averaged

for tr_idx, val_idx in kf.split(stack_train):

    # Training part of stacking features
    X_tr, X_val = stack_train.iloc[tr_idx], stack_train.iloc[val_idx]
    y_tr, y_val = y.iloc[tr_idx], y.iloc[val_idx]

    # Meta-model (Lasso works incredibly well here)
    meta_model = Lasso(alpha=0.0005, max_iter=50000, random_state=42)
    meta_model.fit(X_tr, y_tr)

    # OOF prediction for validation fold
    oof_meta[val_idx] = meta_model.predict(X_val)

    # Prediction for test set (accumulated)
    test_meta += meta_model.predict(stack_test) / kf.get_n_splits()

# Print meta-model performance
print("Meta-model OOF log-RMSE:", rmsle(y, oof_meta))

# Final predictions from stacked model
final_log_preds = test_meta
final_preds = np.expm1(final_log_preds)
final_preds = np.clip(final_preds, 0, None)

# Save submission
sub = test[['Id']].copy()
sub['SalePrice'] = final_preds
sub.to_csv('submission.csv', index=False)

print("\n✓ Final stacked submission saved as submission.csv")

```

Meta-model OOF log-RMSE: 0.11626573157737813

✓ Final stacked submission saved as submission.csv

Model accuracy (Interpreting the final score)

My final stacked meta-model achieved an OOF log-RMSE of **0.1163**.

Hence, this score translates to an approximate **11.6% relative error** in predicting the `SalePrice` of houses.

This means that, on average, my predictions differ from the actual prices by roughly **±11–12%**, which is considered a strong performance for this competition.

Insights:

- Log predictions were converted back to actual prices
- Negative values were clipped
- The final CSV (**submission.csv**) was created for Kaggle submission



Key Learnings

- Feature engineering (TotalSF, Age, baths, porches, TE) adds strong predictive signal.
- Leak-free target encoding prevents data leakage and improves model stability.
- Rebuilding preprocessing after FE keeps train/test feature spaces fully aligned.
- OOF predictions create clean inputs for stacking without overfitting.
- Boosting models capture non-linear patterns that linear models miss.
- Stacking blends strengths of all models and reduces variance.
- Lasso meta-model learns the best combination of base model predictions.